

PROGINF/FTRACE User's Guide (G2AT03E)

SX-Aurora TSUBASA

Proprietary Notice

The information disclosed in this document is the property of NEC Corporation (NEC) and/or its licensors. NEC and/or its licensors, as appropriate, reserve all patent, copyright and other proprietary rights to this document, including all design, manufacturing, reproduction, use and sales rights thereto, except to the extent said rights are expressly granted to others. The information in this document is subject to change at any time, without notice. All product, brand, or trade names in this publication are the trademarks or registered trademarks of their respective owners

Introduction

This PROGINF/FTRACE User's Guide describes the features and usage of the performance analysis tool "PROGINF" and "FTRACE" for the SX-Aurora TSUBASA system.

Targeted readers in this document

This document is written mainly for general users and programmers and it assumes that the reader knows Fortran Compiler (nfort), C Compiler (ncc), C++ Compiler (nc++) and NEC MPI.

How to read this document

This document consists of the following chapters. Please refer to the target reader on the right side of the table and read on.

Chapter	Title	Contents	Targeted Readers
Chapter 1	PROGINF	This chapter explains usage of PROGINF.	general users programmers
Chapter 2	FTRACE	This chapter explains the usage of FTRACE	general users programmers

Related Documents

The following documents have additional information related to PROGINF and FTRACE.

- Fortran Compiler (nfort)
Fortran User's Guide
- C Compiler (ncc) and C++ Compiler (nc++)
C/C++ User's Guide
- NEC MPI
NEC MPI User's Guide

Conventions

The following conventions are used throughout this document.

- Bold text that starts with hyphen (-)
It indicates a command option.
- Dollar sign (\$) in an example of executing command
It indicates **bash** (Bourne-Again Shell) prompt unless otherwise noted.
- **ncc/nc++/nfort/mpincc/mpinc++/mpinfort**
In this document, these commands to compile and link a program are called "compiler".

Contents

Chapter 1	PROGINF	1
1.1	Usage.....	1
1.2	Options	7
1.2.1	Compiler Option.....	7
1.2.2	Environment Variable	7
1.3	Notes.....	8
Chapter 2	FTRACE	9
2.1	Usage.....	9
2.2	User Specified Region	18
2.3	Options	20
2.3.1	Compiler Option.....	20
2.3.2	Environment Variable at the runtime of executable file.....	20
2.3.3	ftrace Command Option	20
2.4	Notes.....	21
Appendix A	Revision History.....	23

Chapter 1 PROGINF

PROGINF provides program execution analysis information throughout the execution of program. This feature is available by linking the program using NEC Compiler.

1.1 Usage

Link a program against PROGINF library. PROGINF library is linked as default, so you don't need to specify any options.

Example When linking C program using ncc.

```
$ ncc source.c
```

Example When linking C++ program using nc++.

```
$ nc++ source.cc
```

Example When linking Fortran program using nfort.

```
$ nfort source.f90
```

Set the environment variable **VE_PROGINF** to **YES** or **DETAIL** and run the executable file.

Example For bash and its variants

```
$ export VE_PROGINF=DETAIL
$ /opt/nec/ve/bin/ve_exec ./a.out
```

Example For csh and its variants

```
$ setenv VE_PROGINF DETAIL
$ /opt/nec/ve/bin/ve_exec ./a.out
```

At the end of program execution, the program execution analysis information is output to standard error (stderr) as follows.

***** Program Information *****		
Real Time (sec)	:	33.060014
User Time (sec)	:	109.886525
Vector Time (sec)	:	104.243699
Inst. Count	:	77145324923
V. Inst. Count	:	16268966755
V. Element Count	:	4164804790539
V. Load Element Count	:	1921282387058
FLOP Count	:	1281280040326
MOPS	:	42021.575083
MOPS (Real)	:	137499.248243
MFLOPS	:	11844.518710
MFLOPS (Real)	:	38756.577193
A. V. Length	:	255.996884
V. Op. Ratio (%)	:	98.660787
L1 Cache Miss (sec)	:	0.414586
CPU Port Conf. (sec)	:	0.000000
V. Arith. Exec. (sec)	:	35.594797
V. Load Exec. (sec)	:	68.634671
VLD LLC Hit Element Ratio (%)	:	0.001852
FMA Element Count	:	2002000063
Power Throttling (sec)	:	0.000000
Thermal Throttling (sec)	:	0.000000
Max Active Threads	:	4
Available CPU Cores	:	8
Average CPU Cores Used	:	3.323850
Memory Size Used (MB)	:	7884.000000
Start Time (date)	:	Mon Feb 19 04:43:34 2018 JST
End Time (date)	:	Mon Feb 19 04:44:08 2018 JST

The meanings of items are as follows.

(*1) output only when VE_PROGINF is set to DETAIL

(*2) output only when VE_PROGINF is set to DETAIL and multi-thread execution

Item	Unit	Description
Real Time	second	Elapsed time
User Time	second	User time
Vector Time	second	Vector instruction execution Time
Inst. Count	-	Number of all instruction executions
V. Inst. Count	-	Number of vector instruction executions
V. Element Count	-	Number of vector instruction execution elements
V. Load Element Count	-	Number of vector instruction load elements
FLOP Count	-	Number of floating-point data execution elements
MOPS	-	Number of million operations performed per “User Time” second
MOPS (Real)	-	Number of million operations performed per “Real Time” second
MFLOPS	-	Number of million floating-point data execution elements executed per “User Time” second
MFLOPS (Real)	-	Number of million floating-point data execution elements executed per “Real Time” second
A. V. Length	-	Average vector length
V. Op. Ratio	percent	Vector Operation Ratio: The ratio of vector operations to all operations
L1 Cache Miss	second	L1 cache miss time
CPU Port Conf.	second	CPU port conflict time (*1)
V. Arith Exec.	second	Vector arithmetic execution time (*1)
V. Load Exec.	second	Vector load execution time (*1)
VLD LLC Hit Element Ratio	percent	The ratio of the elements loaded from LLC to the total elements loaded by vector instructions.

Item	Unit	Description
FMA Element Count	-	Number of FMA execution elements (*1)
Power Throttling	second	Power throttling time (*1)
Thermal Throttling	second	Thermal throttling time (*1)
Max Active Threads	-	Number of maximum active threads (*2)
Available CPU cores	-	Number of available CPU cores (*2)
Average CPU Cores Used	-	Average of CPU cores used (*2)
Memory Size Used	megabyte	Maximum memory size used by program
Start Time (date)	-	The date when the program is started
End Time (date)	-	The date when the program is terminated

PROGINF outputs the program execution analysis information using Aurora HW performance counter. You can control the use of set of the performance counters by the environment variable VE_PERF_MODE and PROGINF can output items corresponding the set. Above output is the case that VE_PERF_MODE is unset or VE_PERF_MODE is set to VECTOR-OP. In this case, PROGINF outputs items related to vector instructions mainly. The following output is the case that VE_PERRF_MODE is set to VECTOR-MEM. In this case, PROGINF outputs items related to vector and memory access mainly.

***** Program Information *****		
Real Time (sec)	:	41.899895
User Time (sec)	:	167.316072
Vector Time (sec)	:	92.309990
Inst. Count	:	175930159073
V. Inst. Count	:	16278237491
V. Element Count	:	564896066618
V. Load Element Count	:	167991674972
FLOP Count	:	432838957109
MOPS	:	4956.779900
MOPS (Real)	:	19821.081454
MFLOPS	:	2583.413548
MFLOPS (Real)	:	10330.507186
A. V. Length	:	34.702533
V. Op. Ratio (%)	:	80.776073
L1 I-Cache Miss (sec)	:	0.256811
L1 O-Cache Miss (sec)	:	32.683728
L2 Cache Miss (sec)	:	32.884266
FMA Element Count	:	105937492420
Required B/F	:	5.773960
Required Store B/F	:	2.195626
Required Load B/F	:	3.578334
Actual Load B/F	:	0.404243
Power Throttling (sec)	:	0.000000
Thermal Throttling (sec)	:	0.000000
Max Active Threads	:	4
Available CPU Cores	:	8
Average CPU Cores Used	:	3.993234
Memory Size Used (MB)	:	1028.000000
Start Time (date)	:	Sun Dec 15 17:03:50 2019 JST
End Time (date)	:	Sun Dec 15 17:04:32 2019 JST

Comparing that VE_PERF_MODE is unset or VE_PERF_MODE is set to VECTOR-OP, PROGINF outputs the following items instead of L1 Cache Miss, CPU Port Conf., V. Arith Exec., V. Load Exec. and VLD LLC Hit Element Ratio.

(*1) output only when VE_PROGINF=DETAIL

(*2) the value over 100 is truncated.

Item	Unit	Description
L1 I-Cache Miss	second	L1 instruction cache miss time
L1 O-Cache Miss	second	L1 operand cache miss time
L2 Cache Miss	second	L2 cache miss time
Required B/F	-	B/F calculated from bytes specified by load and store instructions (*1) (*2)
Required Store B/F	-	B/F calculated from bytes specified by store instructions (*1) (*2)
Required Load B/F	-	B/F calculated from bytes specified by load instructions (*1) (*2)
Actual V. Load B/F	-	B/F calculated from bytes of actual memory access by vector load instructions (*1) (*2)

1.2 Options

Compiler option and environment variable related to PROGINF are as follows.

1.2.1 Compiler Option

Option	Description
-proginf	Link a program against PROGINF library (default)
-no-proginf	Don't link a program against PROGINF library

1.2.2 Environment Variable

Environment Variable	Description
VE_PROGINF	Control whether the program execution analysis information is output at the end of program execution or not. NO: Don't output (default) YES: Output basic information DETAIL: Output detailed information
VE_PROGINF_USE_SIGNAL	Signal SIGUSR1 is used for collecting performance information of thread. When you use this environment variables, please refer to the note of 1.3(1). NO: Don't use signal YES: Use signal (default)
VE_PERF_MODE	Control the HW performance counter set. PROGINF outputs program execution analysis information corresponding to the set. VECTOR-OP: Select the set related to vector operation mainly.(default) VECTOR-MEM: Select the set related to vector and memory access mainly.

1.3 Notes

- (1) PROGINF uses signal SIGUSR1 for collecting performance information of threads as default. For example, using debugger, you can stop the signal by the environment variable `VE_PROGINF_USE_SIGBAL` is set to NO. In this case, PROGINF terminate worker thread of compiler's automatic parallelization and OpenMP by collecting performance information of those threads instead of signal. When PROGINF does not signal, PROGINF cannot collect performance information of the threads except for the threads of compiler's automatic parallelization and OpenMP and the program execution analysis information may not be displayed correctly.
 - Items to be effective during multi thread execution are not shown
 - Performance information of the threads which PROGINF cannot collect is ignored in each item excluding User Time, Real Time and Memory Size Used.
- (2) If compilation and linking correspond to the following, you need to use the option `-pthread` for the link between program and pthread library. If `-lpthread` is specified instead of the option `-pthread`, the program execution analysis information may not be displayed correctly as in (1).
- (3) When shared libraries created with the compiler option `-fopenmp` or `-mparallel` is linked to an executable file, the program execution analysis information may not be displayed correctly as in (1).
- (4) If you use shared libraries created with the option `-pthread`, you need to specify the option `-pthread` again when you link to an executable file, the program execution analysis information may not be displayed correctly as in (1).
- (5) PROGINF does not support multi-process program except for MPI.

Chapter 2 FTRACE

FTRACE is used to obtain performance information such as the CPU usage and vectorization aspect of each function in a program, as well as user regions.

2.1 Usage

Compile and link a program with **-ftrace** to an executable file for performance measurement.

Example When compiling and linking C program using ncc

```
$ ncc -ftrace source.c
```

Example When compiling and linking C++ program using nc++

```
$ nc++ -ftrace source.cc
```

Example When compiling and linking Fortran program using nfort

```
$ nfort -ftrace source.f90
```

For MPI program, **mpincc**, **mpinc++** and **mpinfort** are used to compile C, C++ and Fortran program, respectively.

Example When compiling and linking MPI C program using mpincc

```
$ mpincc -ftrace source.c
```

Example When compiling and linking MPI C++ program using mpinc++

```
$ mpinc++ -ftrace source.cc
```

Example When compiling and linking MPI Fortran program using mpinfort

```
$ mpinfort -ftrace source.f90
```

Run the executable file to measure performance and collect analysis information files. You can run the executable file in the same way as a normal executable file. At the end of execution, one or more analysis information files are generated in a working directory where the program is executed. In the case of non-MPI program, a single analysis information file is created. In the case of MPI program, analysis information file is created for each MPI process. The file name is given as follows.

- In the case of non-MPI Program
ftrace.out
- In the case of MPI Program
ftrace.out.*univ.rank*

Where *univ* is the universe number and *rank* is the value of rank in MPI_COMM_WORLD.

Example In the case of non-MPI program

```
$ /opt/nec/ve/bin/ve_exec ./a.out
$ ls ftrace.out
ftrace.out
```

Example In the case of MPI program

```
$ mpirun -np 4 /opt/nec/ve/bin/ve_exec ./a.out
$ ls ftrace.out.*
ftrace.out.0.0 ftrace.out.0.1 ftrace.out.0.2 ftrace.out.0.3
```

Type **ftrace** command to read the analysis information files.

Example When outputting an analysis list of non-MPI program

```
$ ftrace -f ftrace.out
```

Example When outputting MPI program

```
$ ftrace -f ftrace.out.*
```

Example When outputting MPI process of universe 0 and rank 0

```
$ ftrace -f ftrace.out.0.0
```

The following analysis list is output to the standard output (stdout).

FTRACE ANALYSIS LIST												

Execution Date : Sat Feb 17 12:44:49 2018 JST												
Total CPU Time : 0:03'24"569 (204.569 sec.)												
FREQUENCY	EXCLUSIVE TIME[sec](%)	AVER.TIME [msec]	MOPS	MFLOPS	V.OP RATIO	AVER. V.LEN	VECTOR TIME	L1CACHE MISS	CPU PORT CONF	VLD LLC HIT	E.%	PROC.NAME
1012	49.093(24.0)	48.511	23317.2	14001.4	96.97	83.2	42.132	5.511	0.000	80.32		funcA
160640	37.475(18.3)	0.233	17874.6	9985.9	95.22	52.2	34.223	1.973	2.166	96.84		funcB
160640	30.515(14.9)	0.190	22141.8	12263.7	95.50	52.8	29.272	0.191	2.544	93.23		funcC
160640	23.434(11.5)	0.146	44919.9	22923.2	97.75	98.5	21.869	0.741	4.590	97.82		funcD
160640	22.462(11.0)	0.140	42924.5	21989.6	97.73	99.4	20.951	1.212	4.590	96.91		funcE
53562928	15.371(7.5)	0.000	1819.0	742.2	0.00	0.0	0.000	1.253	0.000	0.00		funcG
8	14.266(7.0)	1783.201	1077.3	55.7	0.00	0.0	0.000	4.480	0.000	0.00		funcH
642560	5.641(2.8)	0.009	487.7	0.2	46.45	35.1	1.833	1.609	0.007	91.68		funcF
2032	2.477(1.2)	1.219	667.1	0.0	89.97	28.5	2.218	0.041	0.015	70.42		funcI
8	1.971(1.0)	246.398	21586.7	7823.4	96.21	79.6	1.650	0.271	0.000	2.58		funcJ

54851346	204.569(100.0)	0.004	22508.5	12210.7	95.64	76.5	154.524	17.740	13.916	90.29		total

The meanings of items are as follows.

Item	Unit	Description
Execution Date	-	Date when execution terminated
Total CPU Time	-	Total CPU time for each function
PROC.NAME	-	Function/Region Name
FREQUENCY	-	Calling count of a function
EXCLUSIVE TIME	second (percent)	The exclusive CPU time (second), and the ratio of the exclusive CPU time required for executing a function to the CPU time required for executing the whole function

Item	Unit	Description
AVER.TIME	millisecond	The average CPU time (millisecond) required to execute one function
MOPS	-	Number of million operations performed per “EXCLUSIVE TIME” second
MFLOPS	-	Number of million floating-point data execution elements executed per “EXCLUSIVE TIME” second
V.OP RATIO	percent	The ratio of vector operations to all operations
AVER.V.LEN	-	Average vector length
VECTOR TIME	second	Vector instruction execution time
L1CACHE MISS	second	L1 cache miss time
CPU PORT CONF.	second	CPU port conflict time
VLD LLC HIT E.%	percent	The ratio of the elements loaded from LLC to the total elements loaded by vector load instructions.

FTRACE outputs analysis list from Aurora HW performance counters at the runtime of the executable file. You can control the use of set of the performance counters by the environment variable VE_PERF_MODE at the runtime of the executable file and FTRACE can output different information corresponding the set. Above list is the case that VE_PERF_MODE is unset or VE_PERF_MODE is set to VECTOR-OP. In this case, FTRACE outputs items related to vector instructions mainly. The following list is the case that VE_PERRF_MODE is set to VECTOR-MEM. In this case, FTRACE outputs items related to vector and memory access mainly.

```

*-----*
FTRACE ANALYSIS LIST
*-----*

Execution Date : Sun Dec 15 21:51:48 2019 JST
Total CPU Time : 0:00'39"862 (39.862 sec.)

```

FREQUENCY	EXCLUSIVE TIME[sec](%)	...	L1ICACHE MISS	L1OCACHE MISS	L2CACHE MISS	REQ. B/F	REQ.ST B/F	REQ.LD B/F	ACT.VLD B/F	FLOP COUNT	FMA ELEM.	PROC.NAME
15562	17.311(43.4)	...	0.005	4.441	4.442	5.12	1.34	3.78	0.10	16689871512	4486524600	funcA
15562	17.235(43.2)	...	0.009	4.007	4.009	5.12	1.34	3.78	0.01	16689871512	4187422960	funcB
253	1.886(4.7)	...	0.003	0.049	0.051	6.53	1.55	4.97	0.73	27492387186	9261379732	funcC
15562	1.392(3.5)	...	0.010	0.013	0.021	6.77	3.42	3.35	0.14	21847989784	3469579024	funcnD
15562	1.074(2.7)	...	0.006	0.073	0.077	6.85	3.53	3.32	0.13	19036465492	3469579024	funcE
3147416	0.618(1.6)	...	0.000	0.030	0.000	8.39	2.41	5.98	0.00	387132168	0	funcF
2	0.211(0.5)	...	0.000	0.000	0.000	39.45	21.54	17.91	0.00	46735724	0	funcG
2	0.099(0.2)	...	0.005	0.003	0.012	13.41	5.38	8.03	3.82	897306183	299101640	funcH
1	0.020(0.1)	...	0.000	0.002	0.002	50.15	26.41	23.75	0.00	3574991	0	funcI
1	0.005(0.0)	...	0.000	0.000	0.000	4.26	1.04	3.22	0.17	156349803	58463060	funcJ
2	0.003(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcK
31124	0.003(0.0)	...	0.000	0.002	0.005	0.00	0.00	0.00	0.00	0	0	funcL
31124	0.002(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcM
4	0.001(0.0)	...	0.000	0.000	0.000	2.74	0.04	2.70	0.96	14376844	4766560	funcN
1	0.000(0.0)	...	0.000	0.000	0.000	100.00	57.77	87.24	0.00	2292	0	funcO
1	0.000(0.0)	...	0.000	0.000	0.000	100.00	100.00	100.00	0.00	265	0	funcP
1	0.000(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcQ
1	0.000(0.0)	...	0.000	0.000	0.000	100.00	100.00	100.00	0.00	14	0	funcR
4	0.000(0.0)	...	0.000	0.000	0.000	100.00	100.00	100.00	0.00	8	0	funcS
1	0.000(0.0)	...	0.000	0.000	0.000	2.60	0.26	2.34	0.10	685289	225700	funcT
44	0.000(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcU
2	0.000(0.0)	...	0.000	0.000	0.000	100.00	100.00	100.00	0.00	4	0	funcV
2	0.000(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcW
1	0.000(0.0)	...	0.000	0.000	0.000	58.60	50.00	8.60	0.00	20	0	funcX
1	0.000(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcY
2	0.000(0.0)	...	0.000	0.000	0.000	0.00	0.00	0.00	0.00	0	0	funcZ
3272238	39.862(100.0)	...	0.038	8.620	8.620	6.26	2.29	3.97	0.30	103262749091	25237042300	total

(... part is omitted)

Comparing that VE_PERF_MODE is unset or VE_PERF_MODE is set to VECTOR-OP, FTRACE outputs the following items instead of L1CACHE MISS, CPU PORT CONF, VLD LLC HIT E.%.

(*1) the value over 100 is truncated.

Item	Unit	Description
L1ICACHE MISS	second	L1 instruction cache miss time
L1OCACHE MISS	second	L1 operand cache miss time
L2CACHE MISS	second	L2 cache miss time
REQ. B/F	-	B/F calculated from bytes specified by load and store instructions (*1)
REQ. ST B/F	-	B/F calculated from bytes specified by store instructions (*1)
REQ. LD B/F	-	B/F calculated from bytes specified by load instructions (*1)
ACT. VLD B/F	-	B/F calculated from bytes of actual memory access by vector load instructions (*1)
FLOP COUNT	-	Number of floating-point data execution elements
FMA ELEM.	-	Number of FMA execution elements

In the case of MPI Program, additional list (MPI special item per MPI UNIVERSE.RANK) is shown as the following format.

ELAPSED TIME[sec]	COMM.TIME [sec]	IDLE TIME / ELAPSED	IDLE TIME [sec]	AVER.LEN / ELAPSED	COUNT [byte]	TOTAL LEN	PROC.NAME [byte]				
12.444	0.000		0.000		0.0	0	0.0 funcA				
9.420	0.000		0.000		0.0	0	0.0 funcB				
7.946	0.000		0.000		0.0	0	0.0 funcG				
7.688	0.000		0.000		0.0	0	0.0 funcC				
7.372	0.000		0.000		0.0	0	0.0 funcH				
5.897	0.000		0.000		0.0	0	0.0 funcD				
5.653	0.000		0.000		0.0	0	0.0 funcE				
1.699	1.475		0.756		3.1K	642560	1.9G funcF				
1.073	1.054		0.987		1.0M	4064	4.0G funcI				
0.704	0.045		0.045		80.0	4	320.0 funcK				

FREQUENCY	EXCLUSIVE TIME[sec] (%)	AVER.TIME [msec]	MOPS	MFLOPS	V.OP RATIO	AVER. V.LEN	VECTOR TIME	L1CACHE MISS	CPU PORT CONF	VLD LLC HIT	PROC.NAME E. %
1012	49.093(24.0)	48.511	23317.2	14001.4	96.97	83.2	42.132	5.511	0.000	80.32	funcA
253	12.089	47.784	23666.9	14215.9	97.00	83.2	10.431	1.352	0.000	79.40	0.0
253	12.442	49.177	23009.2	13811.8	96.93	83.2	10.617	1.406	0.000	81.26	0.1
253	12.118	47.899	23607.4	14180.5	97.00	83.2	10.463	1.349	0.000	79.36	0.2
253	12.444	49.185	23002.8	13808.2	96.93	83.2	10.622	1.404	0.000	81.26	0.3
...											
54851346	204.569(100.0)	0.004	22508.5	12210.7	95.64	76.5	154.524	17.740	13.916	90.29	total
ELAPSED TIME[sec]	COMM.TIME [sec]	COMM.TIME / ELAPSED	IDLE TIME [sec]	IDLE TIME / ELAPSED	AVER.LEN [byte]	COUNT	TOTAL LEN [byte]	PROC.NAME			
12.444	0.000		0.000		0.0	0	0.0	funcA			
12.090	0.000	0.000	0.000	0.000	0.0	0	0.0	0.0			
12.442	0.000	0.000	0.000	0.000	0.0	0	0.0	0.1			
12.119	0.000	0.000	0.000	0.000	0.0	0	0.0	0.2			
12.444	0.000	0.000	0.000	0.000	0.0	0	0.0	0.3			
...											

The meanings of MPI related items are as follows.

Item	Unit	Description
PROC.NAME	-	Function Name or MPI Process (MPIUNIVERSE.MPIRANK)
ELAPSED TIME	second	Elapsed Time
COMM.TIME	second	MPI communication time. This also includes the time of MPI functions
COMM.TIME / ELAPSED	-	The ratio of COMM.TIME to ELAPSED TIME in a function
IDLE TIME	second	Idle time in MPI communication
IDLE TIME / ELAPSED	-	The ratio of IDLE TIME to ELAPSED TIME in a function
AVER.LEN	byte	The average communication amount per MPI communication (TOTAL LEN / COUNT). The unit is bytes, kilobytes, megabytes, gigabytes, terabytes, or petabytes.
COUNT	-	The MPI communication count
TOTAL LEN	byte	The total MPI communication amount. The unit is bytes, kilobytes, megabytes, gigabytes, terabytes, or petabytes.

For parallelized procedures by automatic parallelization or OpenMP, the information of each thread is also output after the information of the function. The information of each thread is listed as the -thraed0, -thread1, -thread2 ... lines. In case of MPI programs, one analysis information file is input to ftrace command to output “each thread” format.

FREQUENCY	EXCLUSIVE TIME[sec] (%)	AVER.TIME [msec]	MOPS	MFLOPS	V.OP RATIO	AVER. V.LEN	VECTOR TIME	L1CACHE MISS	CPU CONF	PORT HIT	VLD E. %	LLC	PROC.NAME
40160	40.688 (26.8)	1.013	3101.6	0.0	0.00	0.0	0.000	0.262	0.000	0.00	funcE		
40160	40.191 (26.4)	1.001	3105.9	0.0	0.00	0.0	0.000	0.264	0.000	0.00	funcD		
40160	9.340 (6.1)	0.233	17931.1	10016.3	95.20	52.2	8.572	0.435	0.541	95.92	funcB		
40160	8.217 (5.4)	0.205	20558.4	11385.3	95.48	52.8	7.828	0.099	0.636	92.91	funcC		
160640	8.090 (5.3)	0.050	257.6	0.0	31.45	6.9	5.275	0.645	0.002	97.92	funcF		
13390732	7.926 (5.2)	0.001	1167.1	359.9	0.00	0.0	0.000	2.520	0.000	0.00	funcG		
160640	6.176 (4.1)	0.038	42375.0	21745.2	98.02	98.3	5.084	0.603	1.147	98.13	funcD\$1		
40160	1.558 (1.0)	0.039	41992.4	21549.0	98.02	98.3	1.274	0.158	0.287	98.28	-thread0		
40160	1.536 (1.0)	0.038	42609.0	21865.3	98.02	98.3	1.264	0.149	0.287	98.41	-thread1		
40160	1.540 (1.0)	0.038	42495.2	21806.8	98.02	98.3	1.268	0.152	0.287	97.92	-thread2		
40160	1.543 (1.0)	0.038	42408.7	21762.4	98.02	98.3	1.279	0.144	0.287	97.92	-thread3		
1012	5.122 (3.4)	5.061	18482.6	11129.8	96.24	83.5	4.122	0.688	0.000	96.46	funcA\$2		
253	1.259 (0.8)	4.975	18802.5	11322.4	96.24	83.5	1.021	0.161	0.000	96.46	-thread0		
253	1.292 (0.8)	5.106	18320.1	11031.9	96.24	83.5	1.033	0.180	0.000	96.46	-thread1		
253	1.294 (0.9)	5.115	18288.9	11013.1	96.24	83.5	1.035	0.181	0.000	96.46	-thread2		
253	1.277 (0.8)	5.049	18528.0	11157.1	96.24	83.5	1.032	0.167	0.000	96.46	-thread3		
...													
14048364	152.037(100.0)	0.011	9379.8	4107.7	74.34	76.3	46.350	8.509	3.477	90.34	total		
ELAPSED TIME[sec]	COMM.TIME [sec]	COMM.TIME / ELAPSED	IDLE TIME [sec]	IDLE TIME / ELAPSED	AVER.LEN [byte]	COUNT	TOTAL LEN [byte]	PROC.NAME					
50.531	0.000	0.000	0.000	0.000	0.0	0	0.0	funcD					
50.415	0.000	0.000	0.000	0.000	0.0	0	0.0	funcE					
11.919	0.000	0.000	0.000	0.000	0.0	0	0.0	funcG					
9.376	0.000	0.000	0.000	0.000	0.0	0	0.0	funcB					
8.591	0.000	0.000	0.000	0.000	0.0	0	0.0	funcH					
8.254	0.000	0.000	0.000	0.000	0.0	0	0.0	funcC					
8.138	7.782	0.956	5.987	0.736	3.1K	160640	490.2M	funcF					
3.863	2.628	0.680	2.557	0.662	1.0M	1016	1017.1M	funcI					
2.105	0.000	0.000	0.000	0.000	0.0	0	0.0	funcA					
1.570	0.000	0.000	0.000	0.000	0.0	0	0.0	funcD\$1					
1.570	0.000	0.000	0.000	0.000	0.0	0	0.0	-thread0					
1.549	0.000	0.000	0.000	0.000	0.0	0	0.0	-thread1					
1.552	0.000	0.000	0.000	0.000	0.0	0	0.0	-thread2					
1.557	0.000	0.000	0.000	0.000	0.0	0	0.0	-thread3					
...													

2.2 User Specified Region

FTRACE has feature to analyze use specified region. User can specify beginning and end of region like below. The performance information of a user-specified region is listed at the end of the analysis list together with the corresponding ID.

Function Prototype

```
// Beginning of user-specified region
extern int ftrace_region_begin(const char *id);

// End of user-specified region
extern int ftrace_region_end(const char *id);
```

Arguments

Any string (ID) can be specified to distinguish a user-specified region. Each user-specified region should have an identical ID at the beginning and end. The ID is output in the PROC.NAME column of the analysis list. The ID cannot be omitted.

Return values

Name	Description
FTRACE_OK	Normal ending
FTRACE_EINVAL	The argument is invalid
FTRACE_ENOBEGIN	Corresponding beginning of user-specified region does not exist
FTRACE_EMAXNUM	The number of user-specified regions has reached the limit of 200

Program Example

```
#include <ftrace.h>
...
(void) ftrace_region_begin("loop#1"); // outside region begin

for (i = 0; i < n; i++) {
    ...
}

(void) ftrace_region_begin("loop#2"); // inside region begin
for (j = 0; j < n; j++) {
    ...
}
(void) ftrace_region_end("loop#2");    // inside region end

(void) ftrace_region_end("loop#1");    // outside region end
```

Output Example

FREQUENCY	EXCLUSIVE TIME[sec](%)	AVER.TIME [msec]	MOPS	MFLOPS	V.OP RATIO	AVER. V.LEN	VECTOR TIME	L1CACHE MISS	CPU PORT CONF	VLD LLC HIT E.%	PROC.NAME
1012	49.093(24.0)	48.511	23317.2	14001.4	96.97	83.2	42.132	5.511	0.000	80.32	funcA
160640	37.475(18.3)	0.233	17874.6	9985.9	95.22	52.2	34.223	1.973	2.166	96.84	funcB
160640	30.515(14.9)	0.190	22141.8	12263.7	95.50	52.8	29.272	0.191	2.544	93.23	funcC
160640	23.434(11.5)	0.146	44919.9	22923.2	97.75	98.5	21.869	0.741	4.590	97.82	funcD
160640	22.462(11.0)	0.140	42924.5	21989.6	97.73	99.4	20.951	1.212	4.590	96.91	funcE
53562928	15.371(7.5)	0.000	1819.0	742.2	0.00	0.0	0.000	1.253	0.000	0.00	funcG
8	14.266(7.0)	1783.201	1077.3	55.7	0.00	0.0	0.000	4.480	0.000	0.00	funcH
642560	5.641(2.8)	0.009	487.7	0.2	46.45	35.1	1.833	1.609	0.007	91.68	funcF
2032	2.477(1.2)	1.219	667.1	0.0	89.97	28.5	2.218	0.041	0.015	70.42	funcI
8	1.971(1.0)	246.398	21586.7	7823.4	96.21	79.6	1.650	0.271	0.000	2.58	funcJ

54851346	204.569(100.0)	0.004	22508.5	12210.7	95.64	76.5	154.524	17.740	13.916	90.29	total
62248	37.709(18.4)	0.606	2200.2	1026.4	0.00	0.0	0.000	0.532	0.000	20.00	loop#1
2032	4.834(2.4)	2.379	415.8	0.0	28.61	6.3	4.098	0.246	0.000	0.00	loop#2

2.3 Options

Compiler option related to FTRACE and ftrace command option are as follows.

2.3.1 Compiler Option

Option	Description
<code>-ftrace</code>	Create an object file and the executable file for ftrace function.

2.3.2 Environment Variable at the runtime of executable file.

Environment Variable	Description
VE_PERF_MODE	Control the HW performance counter set. FTRACE outputs program execution analysis information corresponding to the set. VECTOR-OP: Select the set related to vector operation mainly.(default) VECTOR-MEM: Select the set related to vector and memory access mainly.

2.3.3 ftrace Command Option

Option	Description
<code>-all</code>	Outputs performance analysis information of all functions in the analysis information files are specified in <code>-f</code> .
<code>-num n</code>	Only <i>n</i> high-cost functions are output. When <code>-all</code> and <code>-num</code> options are specified at the same time, the last option is effective.
<code>-f file-name(s)</code>	Outputs the analysis list for the analysis information file specified by <code>-f</code> . When <code>-f</code> is omitted, outputs the analysis list for the file <code>ftrace.out</code> . Multiple analysis information files can be specified by blanks. Then, the list is output in the following rules: Only 10 high-cost functions are output. You can change the number by <code>-all</code> or <code>-num</code> .

2.4 Notes

- (1) FTRACE does not support multi-threaded program except automatic parallelization and OpenMP parallelization of nfort, ncc or nc++ and multi-process program except MPI.
- (2) When a parallelized function is compiled without -ftrace, the performance information of the threads except master thread is not included in the analysis results.
- (3) The performance information of the threads except master thread created by the parallelized functions of NLC (NEC Numeric Library Collection) is not included in the analysis results.
- (4) The measurement program calls the measurement routines at the entrance and exit of functions or before and after a user-specified region. Therefore, when the measurement program contains a lot of function calls, its execution time may increase dramatically. In this case, the source files that include functions whose performance information is required must be compiled using -ftrace.

Appendix A Revision History

Edition	Issue	Changes
1	February 2018	Original
2	May 2018	<ul style="list-style-type: none">• Added 1.3 Notes• Updated Program Example in 2.3 User Specified Region• Added 2.4 Notes
3	December 2018	<ul style="list-style-type: none">• Added 1.3 Options• Updated 1.4 Notes• Added 2.4 Options• Updated 2.5 Notes
4	May 2019	<ul style="list-style-type: none">• Updated 2.5 Notes
5	January 2020	<ul style="list-style-type: none">• Updated 1.1 Usage• Updated 1.2 Options• Updated 1.3 Notes• Updated 2.1 Usage• Updated 2.3 Options• Updated 2.4 Notes

SX-Aurora TSUBASA system software

PROGINF/FTRACE User's Guide

February	2018	Edition	1
----------	------	---------	---

January	2020	Edition	5
---------	------	---------	---

NEC Corporation

© NEC Corporation 2018-2020

No part of this document may be reproduced, in any form or by any means,
without permission from NEC Corporation.